

```

'''
Youssef Nasser Abdelhafeez      20200674
Toka Hamdy Muhammed Saeed      20201048
Menna Mohammed Ibrahim        20201178
Merahan Soliman Mohammed      20200574
Abdelaziz Ashraf Abdelaziz      20200321
'''

'\n\nYoussef Nasser Abdelhafeez      20200674\n\nToka Hamdy Muhammed
Saeed      20201048\n\nMenna Mohammed Ibrahim        20201178\n\n
Merahan Soliman Mohammed      20200574\n\nAbdelaziz Ashraf
Abdelaziz      20200321\n\n'

# import libraries.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from skimage.transform import resize
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

## (a) Data Exploration and preparation:
# Load the dataset and perform initial data exploration.
data_train = pd.read_csv('/kaggle/input/mnist-in-csv/mnist_train.csv')
data_test = pd.read_csv('/kaggle/input/mnist-in-csv/mnist_test.csv')

# Visualize the distribution of classes (labels)
data_train['label'].value_counts().sort_index().plot(kind='bar')
plt.title('Distribution of Classes')
plt.xlabel('Class (Digit)')
plt.ylabel('Count')
plt.show()

# Begin by familiarizing yourself with the dataset.
print("First five rows of the training data:")
print(data_train.head())

# Identify the number of unique classes.
num_classes = data_train['label'].nunique()
print("Number of unique classes:", num_classes)

# Identify the number of features.

```

```

num_features = len(data_train.columns) - 1 # Subtract 1 for the label
column
print("Number of features:", num_features)

# Check for missing values.
missing_values = data_train.isnull().sum()
print("\nMissing values in the train data:")
print(missing_values)
missing_values1 = data_test.isnull().sum()
print("\nMissing values in the test data:")
print(missing_values1)

# Extract labels and pixel (target and features) values from the
training data.
train_labels = data_train['label'].values # Target
train_pixels = data_train.drop('label', axis=1).values # Features

test_labels = data_test['label'].values # Target
test_pixels = data_test.drop('label', axis=1).values # Features

# Normalize each image by dividing each pixel by 255.
train_pixels_normalized = train_pixels / 255.0
test_pixels_normalized = test_pixels / 255.0

# Resize images to dimensions of 28 by 28.
train_images_resized = np.array([resize(img.reshape(28, 28), (28, 28))
for img in train_pixels_normalized])
test_images_resized = np.array([resize(img.reshape(28, 28), (28, 28))
for img in test_pixels_normalized])

# After resizing, visualize some images to verify the correctness of
the reshaping process.
num_images = 6
plt.figure(figsize=(10, 4))

for i in range(num_images):
    plt.subplot(2, num_images, i + 1)
    plt.imshow(train_pixels_normalized[i].reshape(28, 28),
cmap='gray')
    plt.title(f"Original {i}")

    plt.subplot(2, num_images, i + 1 + num_images)
    plt.imshow(train_images_resized[i], cmap='gray')
    plt.title(f"Resized {i}")

plt.tight_layout()
plt.show()

```


0

	28x21	28x22	28x23	28x24	28x25	28x26	28x27	28x28
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0

[5 rows x 785 columns]

Number of unique classes: 10

Number of features: 784

Missing values in the train data:

label 0

1x1 0

1x2 0

1x3 0

1x4 0

..

28x24 0

28x25 0

28x26 0

28x27 0

28x28 0

Length: 785, dtype: int64

Missing values in the test data:

label 0

1x1 0

1x2 0

1x3 0

1x4 0

..

28x24 0

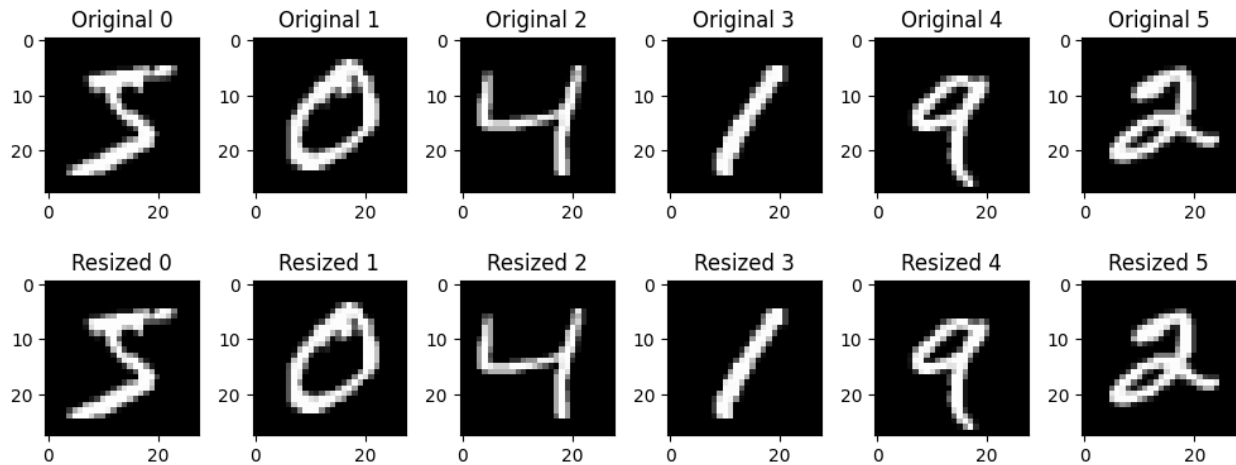
28x25 0

28x26 0

28x27 0

28x28 0

Length: 785, dtype: int64



```
import keras
import seaborn as sns
from keras.layers import Dense
from keras.optimizers import SGD
from keras.models import Sequential
from sklearn.metrics import confusion_matrix

# Handling the shape of the data.
X_train = X_train.reshape(X_train.shape[0],784)
X_test = X_test.reshape(X_test.shape[0],784)
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

# Experiment 1

# Architecture 1
first_model = Sequential()
first_model.add(Dense(units = 512, activation = 'relu', input_shape =
(784,)))
first_model.add(Dense(units = 256, activation = 'relu'))
first_model.add(Dense(units = 128, activation = 'relu'))
first_model.add(Dense(units = 64, activation = 'relu'))
first_model.add(Dense(units = 10, activation = 'softmax'))
first_model.compile(optimizer = SGD(0.001), loss =
'categorical_crossentropy', metrics = ['accuracy'])
first_model.fit(X_train, y_train, batch_size = 32, epochs =
10,verbose = 1)
accuracy1 = first_model.evaluate(x = X_test,y = y_test,batch_size =
32)
print("Score for model 1 =",accuracy1[1])

# Architecture 2
second_model = Sequential()
second_model.add(Dense(units = 256, activation = 'relu', input_shape =
(784,)))
```

```

second_model.add(Dense(units = 64, activation = 'sigmoid'))
second_model.add(Dense(units = 10, activation = 'softmax'))
second_model.compile(optimizer = SGD(0.0001), loss =
'categorical_crossentropy', metrics = ['accuracy'])
second_model.fit(X_train, y_train, batch_size = 64, epochs =
10, verbose = 1)
accuracy2 = second_model.evaluate(x = X_test, y = y_test, batch_size =
64)
print("Score for model 2 =", accuracy2[1])

```

Get the confusion matrix of the best model.

```

def conf_matrix_ann(model):
    print("ANN Confusion Matrix.")
    # Get predictions from the models
    y_pred = np.argmax(model.predict(X_test), axis=1)
    # Calculate the confusion matrices
    conf_matrix = confusion_matrix(np.argmax(y_test, axis=1), y_pred)
    # Display the confusion matrix using seaborn
    plt.figure(figsize = (8, 6))
    sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = 'Blues',
xticklabels = range(10), yticklabels = range(10))
    plt.title('Confusion Matrix for the best architecture')
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.show()

```

```

ann_best_model = Sequential()

```

```

if(accuracy1[1] >= accuracy2[1]):
    # Save the best architecture.
    ann_best_model = first_model
else:
    # Save the best architecture.
    ann_best_model = second_model

```

```

Ann_accuracy = ann_best_model.evaluate(x = X_test, y =
y_test, batch_size = 64)
print("Score for best architecture =", Ann_accuracy[1])

```

Epoch 1/10

```

1407/1407 [=====] - 6s 3ms/step - loss:
1.9513 - accuracy: 0.5033

```

Epoch 2/10

```

1407/1407 [=====] - 5s 3ms/step - loss:
0.9950 - accuracy: 0.7875

```

Epoch 3/10

```

1407/1407 [=====] - 5s 3ms/step - loss:
0.5651 - accuracy: 0.8585

```

Epoch 4/10

```

1407/1407 [=====] - 5s 3ms/step - loss:

```

```
0.4400 - accuracy: 0.8813
Epoch 5/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.3825 - accuracy: 0.8939
Epoch 6/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.3477 - accuracy: 0.9026
Epoch 7/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.3226 - accuracy: 0.9096
Epoch 8/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.3031 - accuracy: 0.9142
Epoch 9/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.2871 - accuracy: 0.9189
Epoch 10/10
1407/1407 [=====] - 5s 3ms/step - loss:
0.2734 - accuracy: 0.9232
469/469 [=====] - 1s 2ms/step - loss: 0.2770
- accuracy: 0.9206
Score for model 1 = 0.9205999970436096
Epoch 1/10
704/704 [=====] - 3s 3ms/step - loss: 2.6129
- accuracy: 0.1040
Epoch 2/10
704/704 [=====] - 2s 3ms/step - loss: 2.5413
- accuracy: 0.1135
Epoch 3/10
704/704 [=====] - 2s 3ms/step - loss: 2.4831
- accuracy: 0.1243
Epoch 4/10
704/704 [=====] - 2s 3ms/step - loss: 2.4356
- accuracy: 0.1371
Epoch 5/10
704/704 [=====] - 2s 3ms/step - loss: 2.3965
- accuracy: 0.1502
Epoch 6/10
704/704 [=====] - 2s 3ms/step - loss: 2.3640
- accuracy: 0.1638
Epoch 7/10
704/704 [=====] - 2s 3ms/step - loss: 2.3367
- accuracy: 0.1767
Epoch 8/10
704/704 [=====] - 2s 3ms/step - loss: 2.3135
- accuracy: 0.1892
Epoch 9/10
704/704 [=====] - 2s 3ms/step - loss: 2.2936
- accuracy: 0.2006
```

```

Epoch 10/10
704/704 [=====] - 2s 3ms/step - loss: 2.2763
- accuracy: 0.2107
235/235 [=====] - 1s 2ms/step - loss: 2.2721
- accuracy: 0.2187
Score for model 2 = 0.21873334050178528
235/235 [=====] - 1s 3ms/step - loss: 0.2770
- accuracy: 0.9206
Score for best architecture = 0.9205999970436096

```

```

from sklearn.model_selection import GridSearchCV, cross_val_predict
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, ConfusionMatrixDisplay

# Split the training data (mnist_train) into training and validation
sets.
X_train_flat =
train_images_resized.reshape(train_images_resized.shape[0], -1)
X_train, X_test, y_train, y_test = train_test_split(X_train_flat,
train_labels, test_size=0.25, random_state=42)

# Trying three different values for neighbors.
trials_of_grid={
    'n_neighbors': [3,5,7],
    'weights': ['uniform', 'distance']
}

KNN_Model=KNeighborsClassifier()

# KNN Model using Grid Search
KNN_Grid_Search=GridSearchCV(KNN_Model, trials_of_grid, cv=5)
KNN_Grid_Search.fit(X_train, y_train)

# finding best parameters for nearest neighbors and weight too
Best_K=KNN_Grid_Search.best_params_
Best_KNN_Model=KNeighborsClassifier(n_neighbors=Best_K['n_neighbors'],
weights=Best_K['weights'])

Best_KNN_Model.fit(X_train, y_train)
y_predict=Best_KNN_Model.predict(X_test)

# calculating accuracy
Knn_accuracy=accuracy_score(y_predict, y_test)
print("KNN-Accuracy =", Knn_accuracy)

# Confusion Matrix
def conf_matrix_knn():
    print("KNN Confusion Matrix.")

```



```

predications_KNN=cross_val_predict(Best_KNN_Model,X_train,y_train,cv=5
)
Confusion_Matrix=confusion_matrix(y_train,predications_KNN)
# Display Confusion Matrix
cm_fig, cm_ax = plt.subplots(figsize=(10, 10))
Confusion_Matrix_display=ConfusionMatrixDisplay(Confusion_Matrix)
Confusion_Matrix_display.plot(ax=cm_ax, cmap=plt.cm.RdPu)
plt.show()

```

KNN-Accuracy = 0.9716

```
import joblib
```

```
# Save the best model, then reload it in a separate file.
```

```

if(Ann_accuracy[1] > Knn_accuracy):
    print("ANN is better.")
    # Get the confusion matrix.
    conf_matrix_ann(ann_best_model)
    ann_best_model.save('best_model.h5')
    best_model_loaded = keras.models.load_model('best_model.h5')

```

```
else:
```

```

    print("KNN is better.")
    # Get the confusion matrix.
    conf_matrix_knn()
    joblib.dump(Best_KNN_Model, 'best_knn_model.joblib')
    best_model_loaded = joblib.load('best_knn_model.joblib')

```


```

y_pred = best_model_loaded.predict(test_features)
accuracy=accuracy_score(y_pred, test_labels)
print("The accuracy for the best model loaded =", accuracy)

```

KNN is better.

KNN Confusion Matrix.



The accuracy for the best model loaded = 0.9692